

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: # Import churn.csv
telco = pd.read_csv('../data/telco_data.csv')
telco
```

```
Out[3]:
```

	Account_Length	Vmail_Message	Day_Mins	Eve_Mins	Night_Mins	Intl_Mins
0	128	25	265.1	197.4	244.7	10.0
1	107	26	161.6	195.5	254.4	13.7
2	137	0	243.4	121.2	162.6	12.2
3	84	0	299.4	61.9	196.9	6.6
4	75	0	166.7	148.3	186.9	10.1
...
3328	192	36	156.2	215.5	279.1	9.9
3329	68	0	231.1	153.4	191.3	9.6
3330	28	0	180.8	288.8	191.9	14.1
3331	184	0	213.8	159.6	139.2	5.0
3332	74	25	234.4	265.9	241.4	13.7

3333 rows x 21 columns

EDA

```
In [465... telco.head()
```

Out [465...

	Account_Length	Vmail_Message	Day_Mins	Eve_Mins	Night_Mins	Intl_Mins	Cus
0	128	25	265.1	197.4	244.7	10.0	
1	107	26	161.6	195.5	254.4	13.7	
2	137	0	243.4	121.2	162.6	12.2	
3	84	0	299.4	61.9	196.9	6.6	
4	75	0	166.7	148.3	186.9	10.1	

5 rows × 21 columns

In [471...

```
telco.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Account_Length        3333 non-null   int64
1   Vmail_Message         3333 non-null   int64
2   Day_Mins              3333 non-null   float64
3   Eve_Mins              3333 non-null   float64
4   Night_Mins            3333 non-null   float64
5   Intl_Mins             3333 non-null   float64
6   CustServ_Calls        3333 non-null   int64
7   Churn                 3333 non-null   int64
8   Intl_Plan             3333 non-null   int64
9   Vmail_Plan            3333 non-null   int64
10  Day_Calls             3333 non-null   int64
11  Day_Charge            3333 non-null   float64
12  Eve_Calls            3333 non-null   int64
13  Eve_Charge           3333 non-null   float64
14  Night_Calls          3333 non-null   int64
15  Night_Charge         3333 non-null   float64
16  Intl_Calls           3333 non-null   int64
17  Intl_Charge          3333 non-null   float64
18  State                3333 non-null   object
19  Area_Code            3333 non-null   int64
20  Phone                3333 non-null   object
dtypes: float64(8), int64(11), object(2)
memory usage: 546.9+ KB
```

In [467...

```
telco.describe()
```

Out [467...

	Account_Length	Vmail_Message	Day_Mins	Eve_Mins	Night_Mins	
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	33
mean	101.064806	8.099010	179.775098	200.980348	200.872037	
std	39.822106	13.688365	54.467389	50.713844	50.573847	
min	1.000000	0.000000	0.000000	0.000000	23.200000	
25%	74.000000	0.000000	143.700000	166.600000	167.000000	
50%	101.000000	0.000000	179.400000	201.400000	201.200000	
75%	127.000000	20.000000	216.400000	235.300000	235.300000	
max	243.000000	51.000000	350.800000	363.700000	395.000000	

In [427... `telco['Churn'].value_counts()`

Out [427... Churn
no 2850
yes 483
Name: count, dtype: int64

In [428... *# Group telco by 'Churn' and compute the mean*
`telco.groupby(['Churn']).mean(numeric_only=True)`

Adapt your code to compute the standard deviation
`telco.groupby(['Churn']).std(numeric_only=True)`

Out [428...

	Account_Length	Vmail_Message	Day_Mins	Eve_Mins	Night_Mins	Intl_Mins
Churn						
no	39.88235	13.913125	50.181655	50.292175	51.105032	2.784489
yes	39.46782	11.860138	68.997792	51.728910	47.132825	2.793190


In [429... *# Count the number of churners and non-churners by State*
`telco.groupby('State')['Churn'].value_counts()`

Filter the data to show only California
`telco[telco['State'] == 'CA'].groupby('State')['Churn'].value_counts()`

Out [429... State Churn
CA no 25
yes 9
Name: count, dtype: int64


In [430... *# Visualize the distribution of 'Intl_Mins'*
`sns.histplot(telco['Intl_Mins'])`

Display the plot
`plt.show()`



```
In [431... # Add "Intl_Plan" as a third variable
sns.boxplot(x = 'Churn',
            y = 'CustServ_Calls',
            data = telco,
            fliersize = 0, # Remove outliers
            # hue = "Intl_Plan" # Add third variable
            )

# Display the plot
plt.show()
```



Data Preprocessing

Encoding Binary Features

```
In [468... pd.set_option('future.no_silent_downcasting', True)

# Replace 'no' with 0 and 'yes' with 1 in 'Vmail_Plan'
# telco['Vmail_Plan'] = telco['Vmail_Plan'].replace({'no': 0 , 'yes': 1})
telco['Vmail_Plan'] = telco['Vmail_Plan'].replace({'no': 0 , 'yes': 1}).infer_objects()

# Replace 'no' with 0 and 'yes' with 1 in 'Churn'
# telco['Churn'] = telco['Churn'].replace({'no': 0 , 'yes': 1})
telco['Churn'] = telco['Churn'].replace({'no': 0 , 'yes': 1}).infer_objects()

# Print the results to verify
print(telco['Vmail_Plan'].head())
print(telco['Churn'].head())
```

```

0    1
1    1
2    0
3    0
4    0
Name: Vmail_Plan, dtype: int64
0    0
1    0
2    0
3    0
4    0
Name: Churn, dtype: int64

```

Encoding Categorical Features using One-Hot Encoding

```
In [469... telco_state = pd.get_dummies(telco, columns=['State'], drop_first=True)
telco_state
```

```
Out[469...      Account_Length  Vmail_Message  Day_Mins  Eve_Mins  Night_Mins  Intl_Mins
0                128                25    265.1    197.4    244.7        10.0
1                107                26    161.6    195.5    254.4        13.7
2                137                 0    243.4    121.2    162.6        12.2
3                 84                 0    299.4     61.9    196.9         6.6
4                 75                 0    166.7    148.3    186.9        10.1
...              ...              ...      ...      ...      ...        ...
3328              192                36    156.2    215.5    279.1         9.9
3329               68                 0    231.1    153.4    191.3         9.6
3330               28                 0    180.8    288.8    191.9        14.1
3331              184                 0    213.8    159.6    139.2         5.0
3332               74                25    234.4    265.9    241.4        13.7
```

3333 rows x 70 columns

```
In [470... # Encode categorical columns
telco['Intl_Plan'] = LabelEncoder().fit_transform(telco['Intl_Plan'])
```

Feature Scaling using Standardization

```
In [472... # Import StandardScaler
from sklearn.preprocessing import StandardScaler

# Select only the numeric columns for scaling
numeric_cols = telco.select_dtypes(include=['float64', 'int64']).columns

# Scale telco
telco_scaled = StandardScaler().fit_transform(telco[numeric_cols])
```

```
# Add column names back for readability
telco_scaled_df = pd.DataFrame(telco_scaled, columns=numeric_cols)

# Print summary statistics
print(telco_scaled_df.describe())
```

	Account_Length	Vmail_Message	Day_Mins	Eve_Mins	\
count	3.333000e+03	3.333000e+03	3.333000e+03	3.333000e+03	
mean	1.470971e-16	7.035077e-17	7.312216e-16	-6.821892e-17	
std	1.000150e+00	1.000150e+00	1.000150e+00	1.000150e+00	
min	-2.513172e+00	-5.917599e-01	-3.301096e+00	-3.963622e+00	
25%	-6.797448e-01	-5.917599e-01	-6.624241e-01	-6.780300e-01	
50%	-1.627644e-03	-5.917599e-01	-6.887677e-03	8.276141e-03	
75%	6.513740e-01	8.695542e-01	6.725198e-01	6.768330e-01	
max	3.564766e+00	3.134591e+00	3.140422e+00	3.209066e+00	

	Night_Mins	Intl_Mins	CustServ_Calls	Churn	Intl_Plan
count	3.333000e+03	3.333000e+03	3.333000e+03	3.333000e+03	3.333000e+03
mean	7.887813e-17	-3.336332e-16	8.527366e-18	5.542788e-17	-4.796643e-17
std	1.000150e+00	1.000150e+00	1.000150e+00	1.000150e+00	1.000150e+00
min	-3.513648e+00	-3.667413e+00	-1.188218e+00	-4.116718e-01	-3.275805e-01
25%	-6.698545e-01	-6.223690e-01	-4.279320e-01	-4.116718e-01	-3.275805e-01
50%	6.485803e-03	2.246393e-02	-4.279320e-01	-4.116718e-01	-3.275805e-01
75%	6.808485e-01	6.672969e-01	3.323545e-01	-4.116718e-01	-3.275805e-01
max	3.839081e+00	3.497397e+00	5.654360e+00	2.429119e+00	3.052685e+00

	Vmail_Plan	Day_Calls	Day_Charge	Eve_Calls	Eve_Charge
count	3.333000e+03	3.333000e+03	3.333000e+03	3.333000e+03	3.333000e+03
mean	-6.608708e-17	-1.934646e-16	-2.835349e-16	3.288365e-16	1.385697e-16
std	1.000150e+00	1.000150e+00	1.000150e+00	1.000150e+00	1.000150e+00
min	-6.183963e-01	-5.005247e+00	-3.301162e+00	-5.025911e+00	-3.963679e+00
25%	-6.183963e-01	-6.695701e-01	-6.623760e-01	-6.583610e-01	-6.783123e-01
50%	-6.183963e-01	2.812491e-02	-6.730063e-03	-5.738630e-03	8.459274e-03
75%	1.617086e+00	6.759846e-01	6.726790e-01	6.970854e-01	6.766695e-01
max	1.617086e+00	3.217588e+00	3.140803e+00	3.508382e+00	3.207980e+00

	Night_Calls	Night_Charge	Intl_Calls	Intl_Charge	Area_Code
count	3.333000e+03	3.333000e+03	3.333000e+03	3.333000e+03	3.333000e+03
mean	-4.903235e-17	-4.370275e-17	-8.527366e-18	2.728757e-16	4.221046e-16
std	1.000150e+00	1.000150e+00	1.000150e+00	1.000150e+00	1.000150e+00
min	-3.429870e+00	-3.515366e+00	-1.820289e+00	-3.668210e+00	-6.888343e-01
25%	-6.699340e-01	-6.676792e-01	-6.011951e-01	-6.164341e-01	-6.888343e-01
50%	-5.505089e-03	4.691242e-03	-1.948306e-01	2.045823e-02	-5.236033e-01
75%	6.589239e-01	6.814562e-01	6.178983e-01	6.706192e-01	1.718817e+00
max	3.827739e+00	3.836763e+00	6.307001e+00	3.496829e+00	1.718817e+00

```
In [473... # Select only the "Intl_Calls" and "Night_Mins" columns
selected_cols = ['Intl_Calls', 'Night_Mins']
telco_selected_scaled_df = telco_scaled_df[selected_cols]
telco_selected_scaled_df.describe()
```

```
Out[473...
           Intl_Calls    Night_Mins
count  3.333000e+03  3.333000e+03
mean   -8.527366e-18   7.887813e-17
std     1.000150e+00   1.000150e+00
min    -1.820289e+00  -3.513648e+00
25%    -6.011951e-01  -6.698545e-01
50%    -1.948306e-01   6.485803e-03
75%     6.178983e-01   6.808485e-01
max     6.307001e+00   3.839081e+00
```

Feature Selection

```
In [474... # Drop the unnecessary features
telco = telco.drop(telco[['Area_Code', 'Phone']], axis=1)

# Verify dropped features
print(telco.columns)
```

```
Index(['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Min
s',
      'Intl_Mins', 'CustServ_Calls', 'Churn', 'Intl_Plan', 'Vmail_Plan',
      'Day_Calls', 'Day_Charge', 'Eve_Calls', 'Eve_Charge', 'Night_Calls',
      'Night_Charge', 'Intl_Calls', 'Intl_Charge', 'State'],
      dtype='object')
```

```
In [479... telco = telco.drop(telco[['State']], axis=1)
```

```
In [475... # Select only numeric columns for correlation computation
numeric_telco = telco.select_dtypes(include=['float64', 'int64'])
numeric_telco.corr() # Identify correlated features
```


Out [475...

	Account_Length	Vmail_Message	Day_Mins	Eve_Mins	Night_Mins
Account_Length	1.000000	-0.004628	0.006216	-0.006757	-0.008955
Vmail_Message	-0.004628	1.000000	0.000778	0.017562	0.007681
Day_Mins	0.006216	0.000778	1.000000	0.007043	0.004323
Eve_Mins	-0.006757	0.017562	0.007043	1.000000	-0.012584
Night_Mins	-0.008955	0.007681	0.004323	-0.012584	1.000000
Intl_Mins	0.009514	0.002856	-0.010155	-0.011035	-0.015207
CustServ_Calls	-0.003796	-0.013263	-0.013423	-0.012985	-0.009288
Churn	0.016541	-0.089728	0.205151	0.092796	0.035493
Intl_Plan	0.024735	0.008745	0.049396	0.019100	-0.028905
Vmail_Plan	0.002918	0.956927	-0.001684	0.021545	0.006079
Day_Calls	0.038470	-0.009548	0.006750	-0.021451	0.022938
Day_Charge	0.006214	0.000776	1.000000	0.007050	0.004324
Eve_Calls	0.019260	-0.005864	0.015769	-0.011430	-0.002093
Eve_Charge	-0.006745	0.017578	0.007029	1.000000	-0.012592
Night_Calls	-0.013176	0.007123	0.022972	0.007586	0.011204
Night_Charge	-0.008960	0.007663	0.004300	-0.012593	0.999999
Intl_Calls	0.020661	0.013957	0.008033	0.002541	-0.012353
Intl_Charge	0.009546	0.002884	-0.010092	-0.011067	-0.015180

In [476...

```
# Compute the correlation matrix
corr_matrix = numeric_telco.corr().abs()

# Select the upper triangle of the correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool), upper=0)
```

Out [476...

	Account_Length	Vmail_Message	Day_Mins	Eve_Mins	Night_Mins
Account_Length	NaN	0.004628	0.006216	0.006757	0.008955
Vmail_Message	NaN	NaN	0.000778	0.017562	0.007681
Day_Mins	NaN	NaN	NaN	0.007043	0.004323
Eve_Mins	NaN	NaN	NaN	NaN	0.012584
Night_Mins	NaN	NaN	NaN	NaN	NaN
Intl_Mins	NaN	NaN	NaN	NaN	NaN
CustServ_Calls	NaN	NaN	NaN	NaN	NaN
Churn	NaN	NaN	NaN	NaN	NaN
Intl_Plan	NaN	NaN	NaN	NaN	NaN
Vmail_Plan	NaN	NaN	NaN	NaN	NaN
Day_Calls	NaN	NaN	NaN	NaN	NaN
Day_Charge	NaN	NaN	NaN	NaN	NaN
Eve_Calls	NaN	NaN	NaN	NaN	NaN
Eve_Charge	NaN	NaN	NaN	NaN	NaN
Night_Calls	NaN	NaN	NaN	NaN	NaN
Night_Charge	NaN	NaN	NaN	NaN	NaN
Intl_Calls	NaN	NaN	NaN	NaN	NaN
Intl_Charge	NaN	NaN	NaN	NaN	NaN

```
In [477... # Find features with correlation greater than a specified threshold (e.g., 0.8)
threshold = 0.8
to_drop = [column for column in upper.columns if any(upper[column] > threshold)]
to_drop
```

Out[477... ['Vmail_Plan', 'Day_Charge', 'Eve_Charge', 'Night_Charge', 'Intl_Charge']

```
In [441... # # Drop the highly correlated features
# numeric_telco_filtered = numeric_telco.drop(columns=to_drop)

# # Display the filtered dataframe
# numeric_telco_filtered.head()
```

```
In [478... # Create the new feature
telco['Avg_Night_Calls'] = telco['Night_Mins'] / telco['Night_Calls']

# Print the first five rows of 'Avg_Night_Calls'
print(telco['Avg_Night_Calls'].head())
```

```
0    2.689011
1    2.469903
2    1.563462
3    2.212360
4    1.544628
Name: Avg_Night_Calls, dtype: float64
```

Churn Prediction

Model Selection

- Logistic Regression: Good baseline, interpretable.
- Random Forests: Captures complex relationships.
- Support Vector Machines (SVMs): Effective with high-dimensional data.

```
In [480... # Define the features
features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins', 'Intl_Mins', 'CustServ_Calls', 'Day_Calls', 'Eve_Calls', 'Night_Calls', 'Intl_Calls', 'Area_Code']

# Define a new customer for prediction
new_customer = pd.DataFrame({
    'Account_Length': [100],
    'Vmail_Message': [20],
    'Day_Mins': [200],
    'Eve_Mins': [150],
    'Night_Mins': [180],
    'Intl_Mins': [10],
    'CustServ_Calls': [2],
    'Day_Calls': [100],
    'Eve_Calls': [100],
    'Night_Calls': [100],
    'Intl_Calls': [5],
    # 'Area_Code': [415]
})
```

```
In [481... # Import LogisticRegression
from sklearn.linear_model import LogisticRegression

# Instantiate the classifier
clf = LogisticRegression()

# Fit the classifier
clf.fit(telco[features], telco['Churn'])

# Predict the label of new_customer
print(clf.predict(new_customer))
```

```
[0]
```

```
/opt/anaconda3/envs/example/lib/python3.12/site-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
In [482... # Import DecisionTreeClassifier
from sklearn.tree import DecisionTreeClassifier

# Instantiate the classifier
clf = DecisionTreeClassifier()

# Fit the classifier
clf.fit(telco[features], telco['Churn'])

# Predict the label of new_customer
print(clf.predict(new_customer))
```

```
[0]
```

Model Evaluation

```
In [483... # Import train_test_split
from sklearn.model_selection import train_test_split

# Create feature variable
X = telco.drop('Churn', axis=1)

# Create target variable
y = telco['Churn']

# Create training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [484... print(X_train.shape)
print(X_test.shape)
print(len(X_test))
```

```
(2333, 18)
(1000, 18)
1000
```

```
In [485... from sklearn.preprocessing import LabelEncoder

# # Encode 'Intl_Plan' column
# le = LabelEncoder()
# X_train['Intl_Plan'] = le.fit_transform(X_train['Intl_Plan'])
# X_test['Intl_Plan'] = le.transform(X_test['Intl_Plan'])

# Import RandomForestClassifier
from sklearn.ensemble import RandomForestClassifier
```

```

# Instantiate the classifier
clf = RandomForestClassifier()

# Fit to the training data
clf.fit(X_train, y_train)

# Compute accuracy
print(clf.score(X_test, y_test))

```

0.955

```

In [486... # Import confusion_matrix
from sklearn.metrics import confusion_matrix

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Print the confusion matrix
print(confusion_matrix(y_test, y_pred))

```

```

[[862   3]
 [ 42  93]]

```

```

In [487... # Create feature variable
X = telco.drop('Churn', axis=1)

# Create target variable
y = telco['Churn']

```

```

In [488... # Import train_test_split
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# Create training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Import RandomForestClassifier
from sklearn.ensemble import RandomForestClassifier

# Instantiate the classifier
clf = RandomForestClassifier()

# Fit to the training data
clf.fit(X_train, y_train)

# Predict the labels of the test set
y_pred = clf.predict(X_test)

# Import confusion_matrix
from sklearn.metrics import confusion_matrix

# Print confusion matrix
print(confusion_matrix(y_test, y_pred))

```

```

[[564   6]
 [ 25  72]]

```

```
In [489... # Import train_test_split
from sklearn.model_selection import train_test_split

# Create training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1)

# Import RandomForestClassifier
from sklearn.ensemble import RandomForestClassifier

# Instantiate the classifier
clf = RandomForestClassifier()

# Fit to the training data
clf.fit(X_train, y_train)

# Predict the labels of the test set
y_pred = clf.predict(X_test)
```

```
In [490... from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

# Print the precision
print(precision_score(y_test, y_pred))

# Print the recall
print(recall_score(y_test, y_pred))
```

```
0.9666666666666667
0.6170212765957447
```


```
In [491... # Generate the probabilities
y_pred_prob = clf.predict_proba(X_test)[:, 1]

# Import roc_curve
from sklearn.metrics import roc_curve

# Calculate the roc metrics
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)

# Plot the ROC curve
plt.plot(fpr, tpr)

# Add labels and diagonal line
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.plot([0, 1], [0, 1], "k--")
plt.show()
```



```
In [456... # Import roc_auc_score
from sklearn.metrics import roc_auc_score

# Print the AUC
print(roc_auc_score(y_test, y_pred_prob))
```

0.9515398550724639

```
In [492... # Instantiate the classifier
clf = RandomForestClassifier()

# Fit to the training data
clf.fit(X_train, y_train)

# Predict the labels of the test set
y_pred = clf.predict(X_test)

# Import f1_score
from sklearn.metrics import f1_score

# Print the F1 score
print(f1_score(y_test, y_pred))
```

0.72

Model Tuning

```
In [493... # Import GridSearchCV
from sklearn.model_selection import GridSearchCV
```

```

# Create the hyperparameter grid
param_grid = {'max_features': ['auto', 'sqrt', 'log2']}

# Call GridSearchCV
grid_search = GridSearchCV(clf, param_grid, cv=3)

# Fit the model
grid_search.fit(X, y)

# Print the optimal parameters
print(grid_search.best_params_)

```

/opt/anaconda3/envs/example/lib/python3.12/site-packages/sklearn/model_selection/_validation.py:540: FitFailedWarning:
3 fits failed out of a total of 9.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

```

-----
3 fits failed with the following error:
Traceback (most recent call last):
  File "/opt/anaconda3/envs/example/lib/python3.12/site-packages/sklearn/model_selection/_validation.py", line 888, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/opt/anaconda3/envs/example/lib/python3.12/site-packages/sklearn/base.py", line 1466, in wrapper
    estimator._validate_params()
  File "/opt/anaconda3/envs/example/lib/python3.12/site-packages/sklearn/base.py", line 666, in _validate_params
    validate_parameter_constraints(
  File "/opt/anaconda3/envs/example/lib/python3.12/site-packages/sklearn/utils/_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto' instead.

warnings.warn(some_fits_failed_message, FitFailedWarning)
/opt/anaconda3/envs/example/lib/python3.12/site-packages/sklearn/model_selection/_search.py:1103: UserWarning: One or more of the test scores are non-finite: [      nan  0.95139514  0.95349535]
  warnings.warn(
{'max_features': 'log2'}

```

In [494...

```

# Import GridSearchCV
from sklearn.model_selection import GridSearchCV

# Create the hyperparameter grid
param_grid = {"max_depth": [3, None],
              "max_features": [1, 3, 10],

```



```

        "bootstrap": [True, False],
        "criterion": ["gini", "entropy"]}]

# Call GridSearchCV
grid_search = GridSearchCV(clf, param_grid, cv=3)

# Fit the model
grid_search.fit(X, y)

# Print the best hyperparameters
print(grid_search.best_params_)

```

```
{'bootstrap': True, 'criterion': 'entropy', 'max_depth': None, 'max_features': 10}
```

```
In [495... # Import RandomizedSearchCV
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

# Create the hyperparameter grid
param_dist = {"max_depth": [3, None],
              "max_features": randint(1, 11),
              "bootstrap": [True, False],
              "criterion": ["gini", "entropy"]}

# Call RandomizedSearchCV
random_search = RandomizedSearchCV(clf, param_dist, cv=3)

# Fit the model
random_search.fit(X, y)

# Print best parameters
print(random_search.best_params_)

```


```
{'bootstrap': True, 'criterion': 'gini', 'max_depth': None, 'max_features': 10}
```

Feature Importance

```
In [496... # Calculate feature importances
importances = clf.feature_importances_

# Create plot
plt.barh(range(X.shape[1]), importances)
plt.show()

```




```
In [497... # Sort importances
sorted_index = np.argsort(importances)

# Create labels
labels = X.columns[sorted_index]

# Clear current plot
plt.clf()

# Create plot
plt.barh(range(X.shape[1]), importances[sorted_index], tick_label=labels)
plt.show()
```



```
In [498... # Import necessary modules
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

# Create training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

# Instantiate the classifier
clf = RandomForestClassifier()

# Fit to the data
clf.fit(X_train, y_train)

# Print the accuracy
print(clf.score(X_test, y_test))
```

0.95

```
In [499... # Import f1_score
from sklearn.metrics import f1_score

# Instantiate the classifier
clf = RandomForestClassifier()

# Fit to the data
clf.fit(X_train, y_train)

# Predict the labels of the test set
y_pred = clf.predict(X_test)

# Print the F1 score
print(f1_score(y_test, y_pred))
```

0.7967479674796748